The background features a large, light gray speedometer with a red needle pointing towards the right. The speedometer's face is partially obscured by a stylized, light gray eye shape. The text "GET ON THE FAST TRACK" is written in a light gray, sans-serif font along the bottom edge of the speedometer. The main title "Module 5.1: Machine Learning Part 1" is centered in a large, black, sans-serif font, and the subtitle "Feature Engineering" is centered below it in a slightly smaller, black, sans-serif font.

# Module 5.1: Machine Learning Part 1

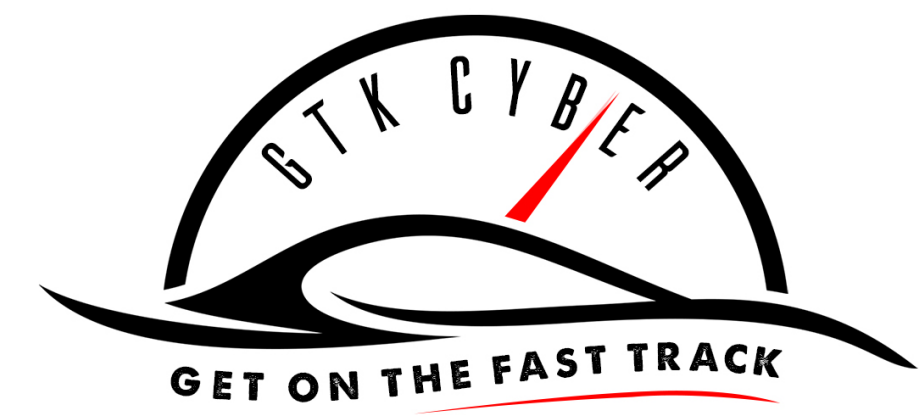
## Feature Engineering

From URL strings to “Features”

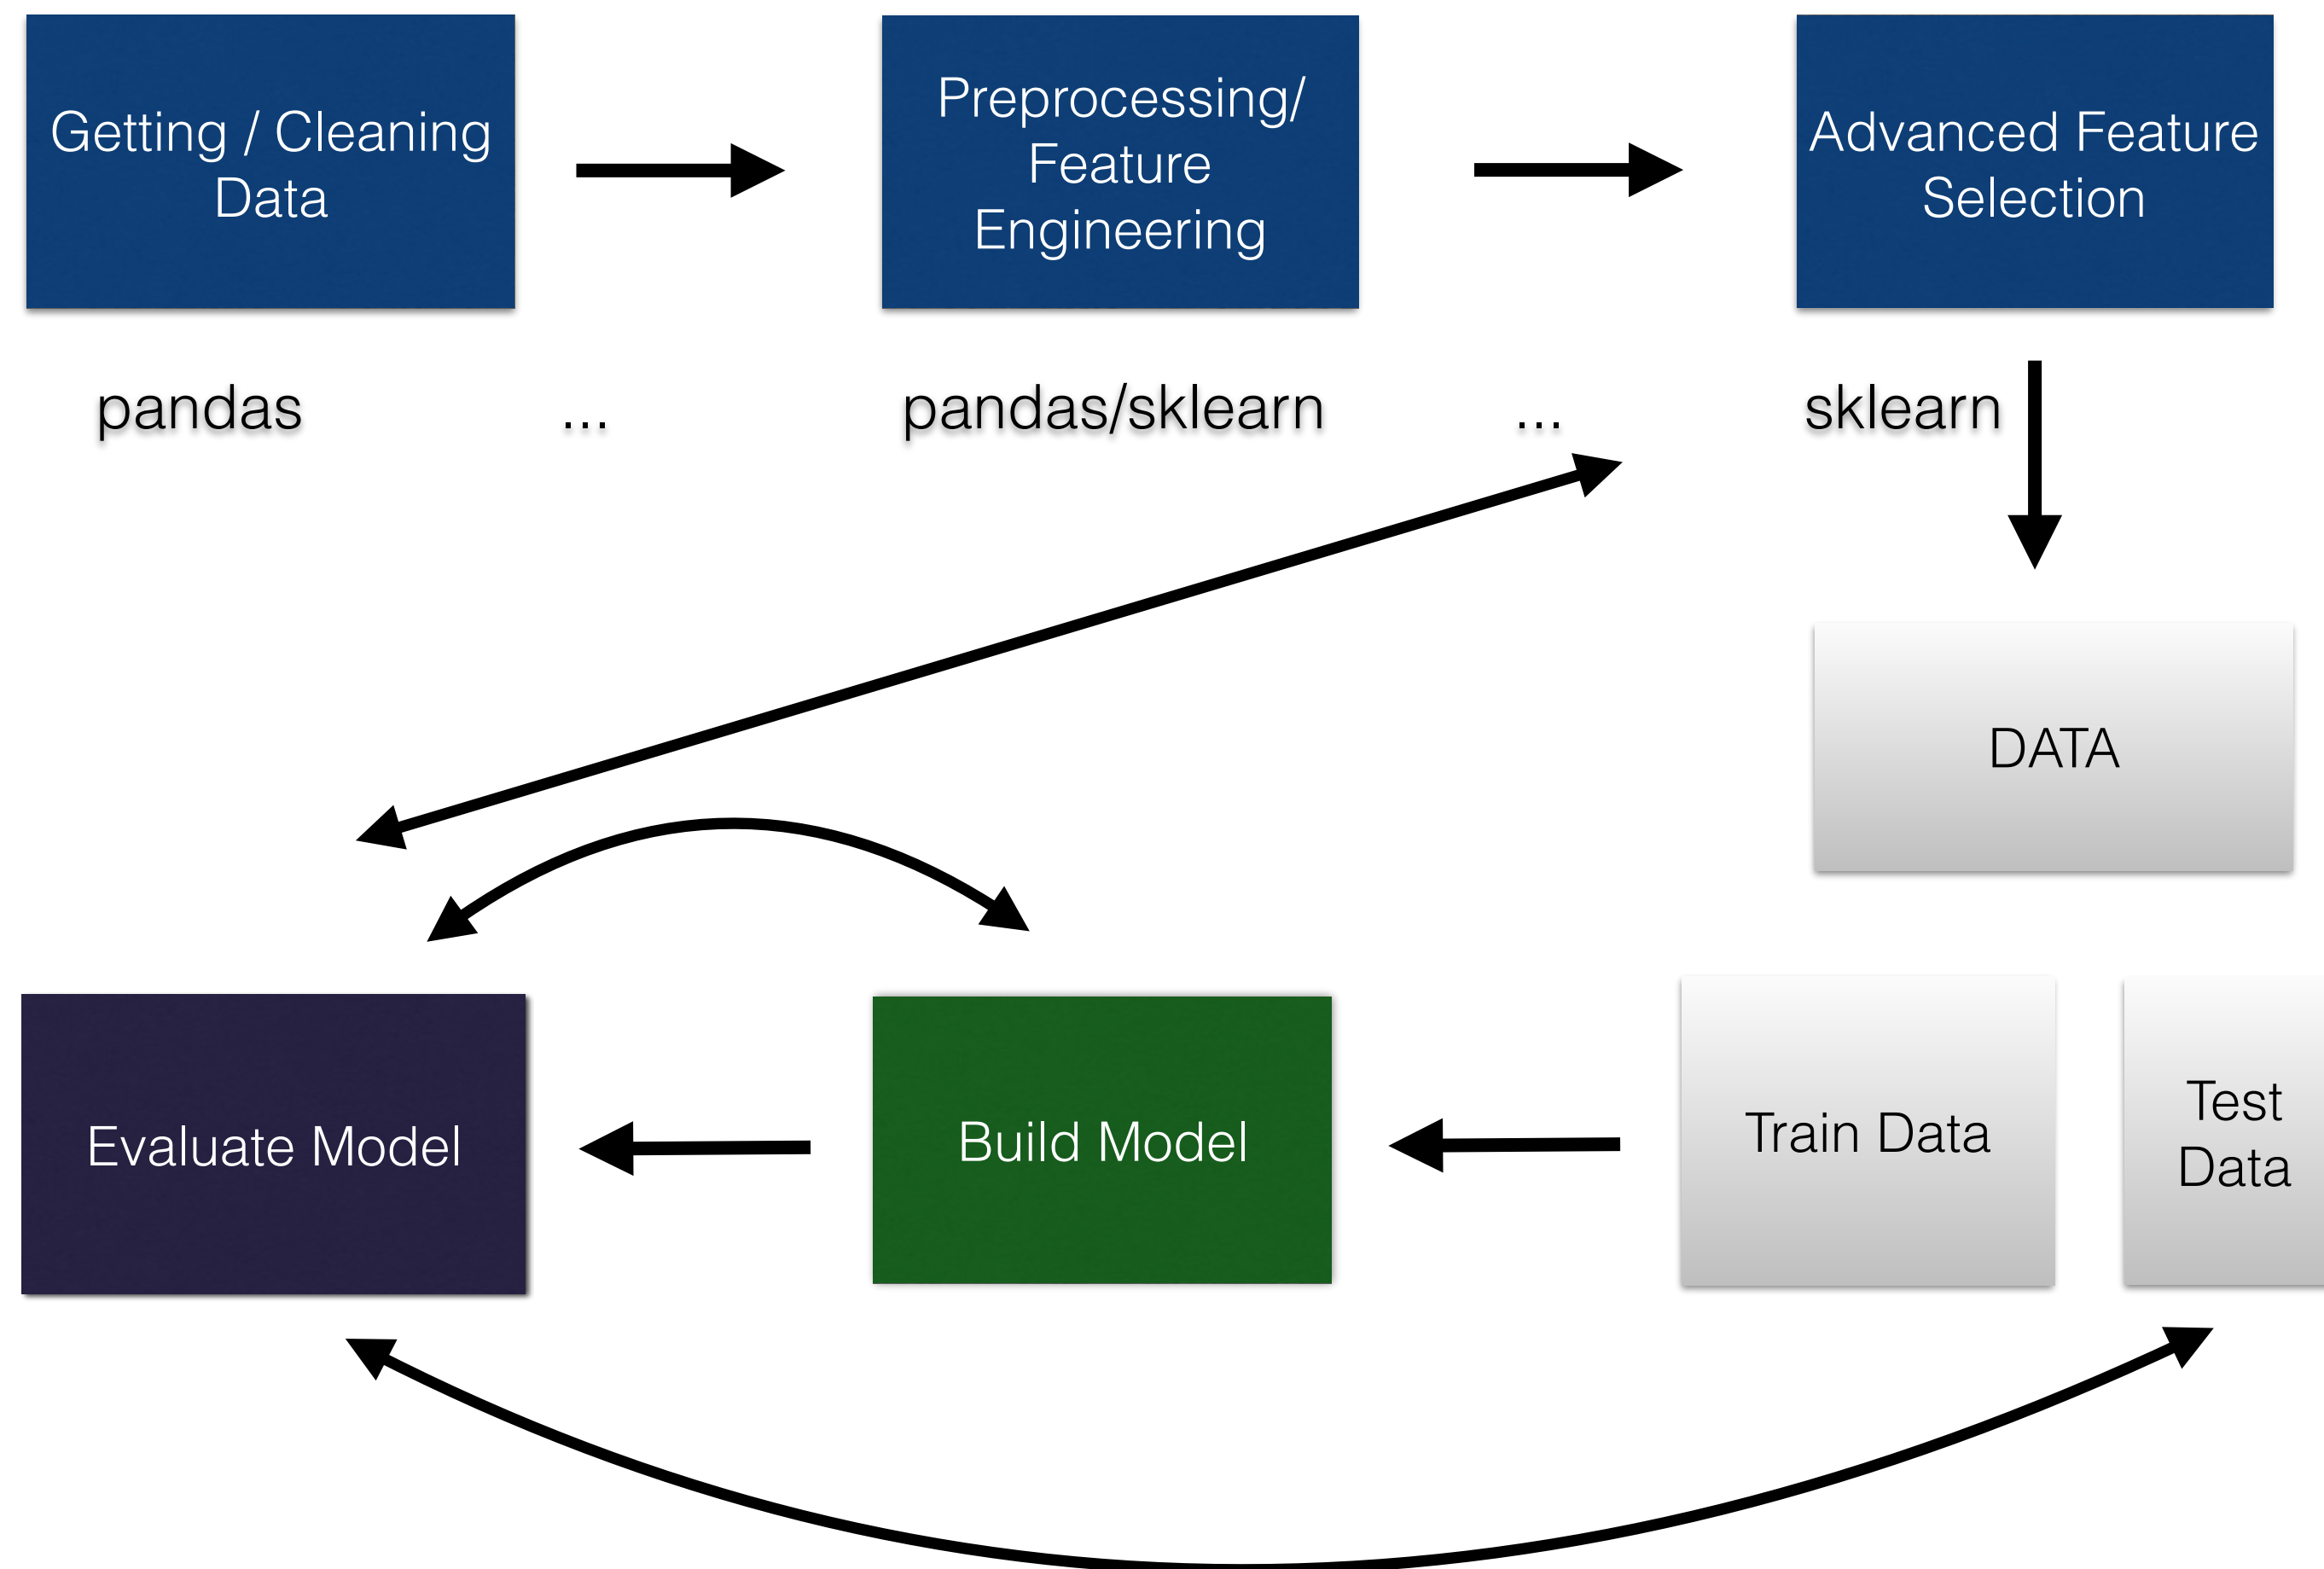


# Agenda

- Feature Selection & Engineering
- Math free overview of classification models
- Evaluating Model Performance
- Improving model performance



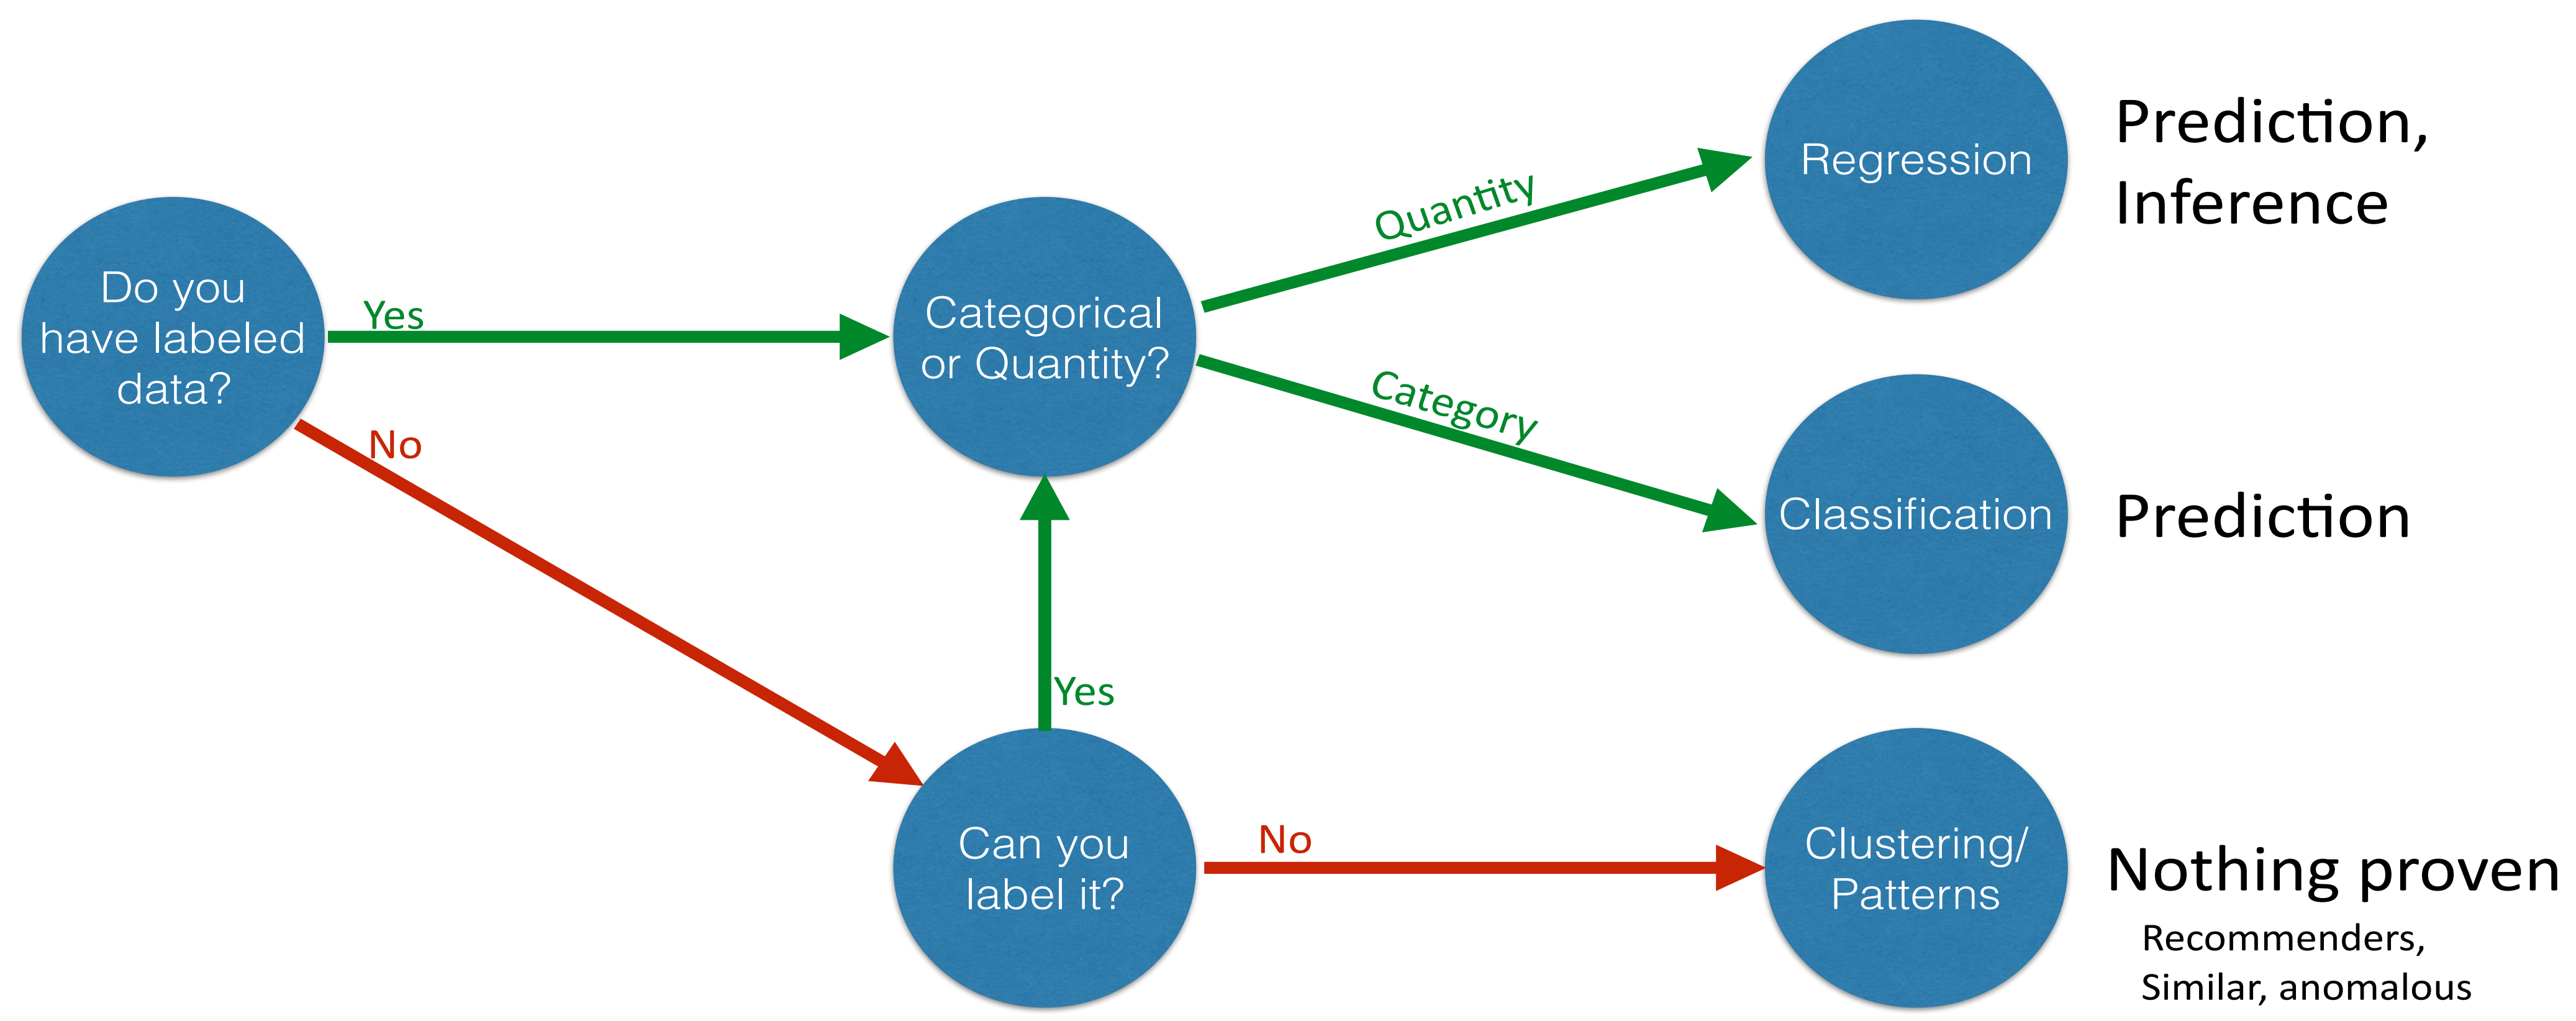
# Machine Learning Process





# Machine Learning Terms

- **Features:** The mathematical representation of the original data. The features are the columns in your data set. Since the features will be a matrix, they are often written as  $X$ .
- **Observations:** The rows of your feature set.
- **Target:** The variable that you are trying to predict. Often represented as  $y$ .







# Features

<http://www.google.com>



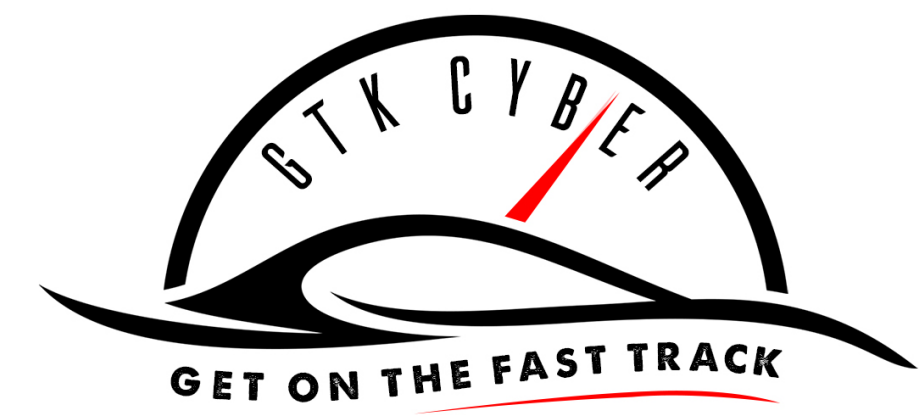


# Features

<http://www.google.com>



domain_length	vowel_count	digit_count
6	3	0



# Representation of URL Knowledge

- Come up with a **representation/set of knowledge** that has **enough complexity** to accurately describe the problem for the computer
- Knowledge here does not mean hard-coded knowledge or formal set of rules
- **The computer rather uses the knowledge we provide to extract patterns and acquire own knowledge**
- We should provide knowledge about reality that has **high variance about the problem** it describes (e.g. a feature that is high when it rains and low when it's sunshine)

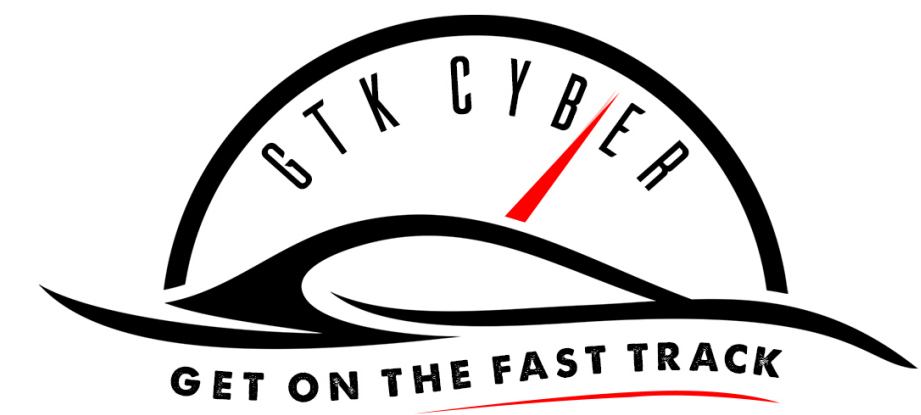




# URL Definition

[https://www.google.com/search?  
q=URL&source=Inms&tbm=isch&sa=X&ved=0ahUKEwjcl6ut-  
IDUAhVEPCYKHdJGDsYQ\\_AUIDCgD&biw=1215&bih=652](https://www.google.com/search?q=URL&source=Inms&tbm=isch&sa=X&ved=0ahUKEwjcl6ut-IDUAhVEPCYKHdJGDsYQ_AUIDCgD&biw=1215&bih=652)

https://	protocol
www	subdomain
google.com	zone apex
google	domain
.com	top-level-domain (tld)
/search?q=URL...	path




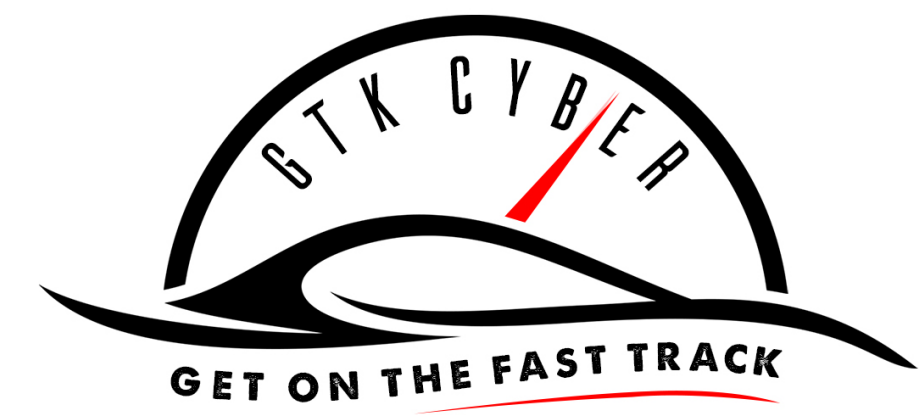
# DNS 101

- **Domain Name Service** (DNS) resolves domain names to IP addresses (like a phone book)
- **Domain Registrars**: authority that signs unique domain names (GoDaddy, BlueGtaor)
- **State of Authority** (SOA): Contains for example name of server for zone, administrator of zone, default time-to-live (ttl = time a DNS record is cached), seconds of secondary name server should wait before checking for updates
- **Root Zone** controlled by Internet Assigned Numbers Authority (IANA)
- **Name Servers** (NS Records): used by tld servers to direct traffic to DNS server (which contains authoritative DNS records)
- **A records** (part of DNS record): “A” stands for IP Address
- **CNAME** (part of DNS record): resolves one domain name to another
- **Autonomous System** (AS) and Border gateway Protocol (BGP) info

Python libraries: `python-whois`, `dnspython`, `tldextract`, `ipaddress`

# DNS Flow





# What makes them different?

## URL BlackList

[amazon-sicherheit.kunden-ueberpruefung.xyz](https://amazon-sicherheit.kunden-ueberpruefung.xyz)

[eclipsehotels.com/language/en-GB/eng.exe](https://eclipsehotels.com/language/en-GB/eng.exe)

[bohicacapital.com/page](https://bohicacapital.com/page)

[summerweb.net](https://summerweb.net)

[ad.getfond.info](https://ad.getfond.info)

[vdula.czystykod.pl/rxdjna2.html](https://vdula.czystykod.pl/rxdjna2.html)

[svision-online.de/mgfi/administrator/  
components/com\\_babackup/classes/fx29id1.txt](https://svision-online.de/mgfi/administrator/components/com_babackup/classes/fx29id1.txt)

## URL WhiteList

[gurufocus.com/stock/PNC](https://gurufocus.com/stock/PNC)

[dvdtalk.ru/review](https://dvdtalk.ru/review)

[333cn.com/zx/zhxw.html](https://333cn.com/zx/zhxw.html)


[made-in-china.com/special/led-lighting](https://made-in-china.com/special/led-lighting)

[google.com/u/0/112261544981697332354/posts](https://google.com/u/0/112261544981697332354/posts)

[youtube.com/watch?v=Qp8MQ4shN6U](https://youtube.com/watch?v=Qp8MQ4shN6U)

[unesco.org/themes/education-sustainable-developm](https://unesco.org/themes/education-sustainable-developm)

[thisisthefirst.com/page/5](https://thisisthefirst.com/page/5)



# Malicious URL Detection Features (Literature)

1. **BlackList Features:** BlackLists suffer from a high false negative rate, but can still be useful as machine learning feature.
2. **Lexical Features:** Capture the property that malicious URLs tend to "look different" from benign URLs. **Contextual information** such as the length of the URL, number of digits, lengths of different parts, entropy of domain name.
3. **Host-based Features:** Properties of web site host. **"Where"** the site is hosted, **"who" owns it** and **"how" it is managed**. API queries are needed (WHOIS, DNS records). Examples: Date of registration, the geolocations, autonomous system (AS) number, connection speed or time-to-live (TTL).
4. **Content-based Features:** Less commonly used feature as it requires **execution of web-page**. Can be not only be not safe, but also increases the computational cost. Examples: HTML or JavaScript based.





# Preparation In Class Exercise

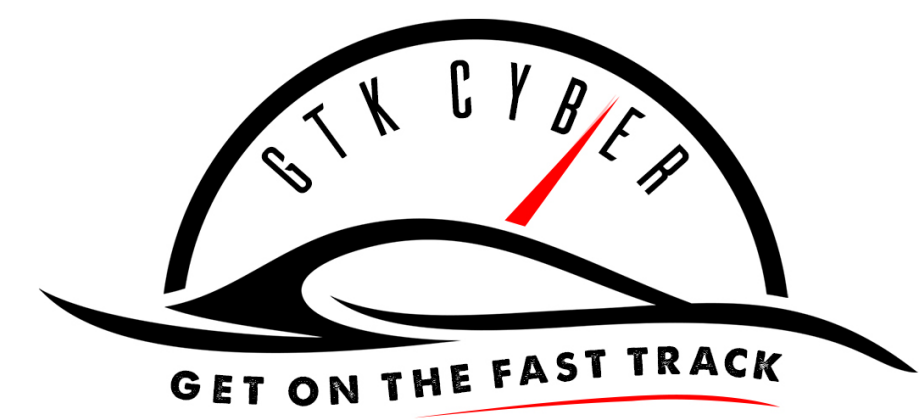
## ML Feature Engineering

### Lexical Features

1. Length of URL
2. Length of domain
3. Count of digits
4. Entropy of domain
5. Position (or index) of the first digit
6. **Bag-of-words** for tld, domain and path parts of the URL

### Host-based Features

1. Time delta between today's date and creation date
2. Check if it is an IP address



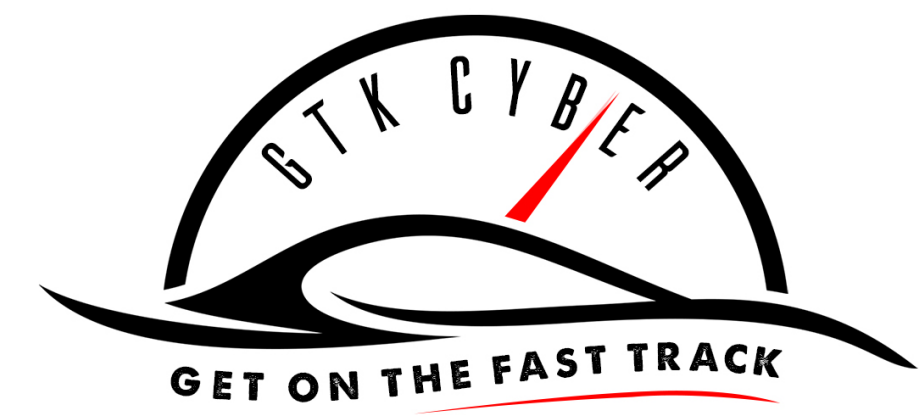
# Data Set (Features and Target)

	url	isMalicious	domain	created
56675	jeita.biz/w/google/drive/document.html?ssl=yes	1	jeita.biz	2012-04-11 17:08:19
73229	sosnovskoe.info/layouts/plugins/mailbox	1	sosnovskoe.info	2011-09-19 09:53:07
60112	teothemes.com/html/mp3pl/blue-preview.jpg	1	teothemes.com	2011-09-08 21:43:00
66946	kfj.cc:162/17852q	1	kfj.cc	2013-08-18 05:52:47
81906	verapdpf.info/db/6d1b281b5c4bbcf3b99228680c232fa	1	verapdpf.info	2016-08-18 07:09:03

Features or X

	isMalicious	isIP	Length	LengthDomain	DigitsCount	EntropyDomain	FirstDigitIndex	com	org	net	...	w	waset
73320	1	0	27	21	0	3.558519	0	0	1	0	...	0	0
30785	0	0	77	11	14	3.095795	22	1	0	0	...	0	0
60789	1	0	141	11	5	3.459432	103	0	1	0	...	0	0
19495	0	0	59	13	20	3.546594	31	1	0	0	...	0	0
45022	1	0	23	11	7	3.277613	13	0	0	0	...	0	0

Target or y



# Bag-of-Words

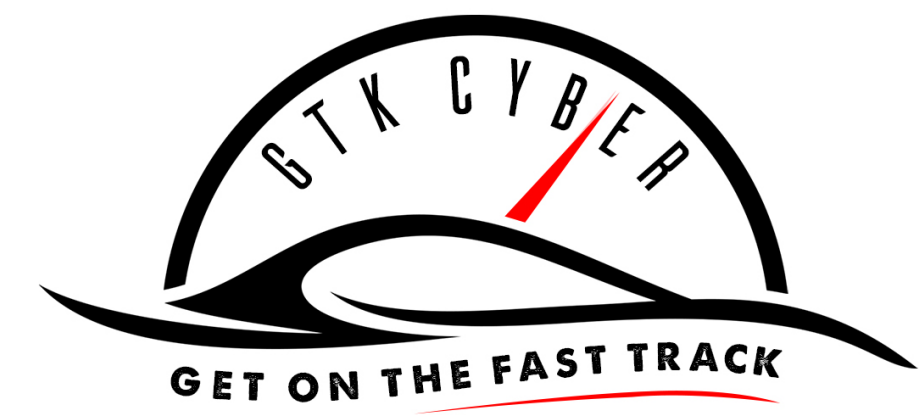
- Bag-of-words model: (Frequency of) occurrence of each word is used as a feature
- Sklearn's `CountVectorizer`: Convert a collection of text documents to a matrix of token counts

Simple Example: Imposing a vocabulary of `top_tlds=[ '.com', '.de', '.uk' ]`  
`CountVectorizer_tlds = CountVectorizer(analyzer='word', vocabulary=top_tlds)`  
`CountVectorizer_tlds = CountVectorizer_tlds.fit(tlds)`  
`matrix_tlds = CountVectorizer_tlds.transform(tlds)`

Bag-of-words model fitting

URL string
...google.ru...
...facebook.com...
...google.de...



‘.com’	‘.de’	‘.uk’
0	0	0
1	0	0
0	1	0



# Preprocessing - an Art Work!

- Imputing missing values
- Scaling/Normalization
- One-Hot Encoding (Encoding categorical features)
- Embedding (e.g. word2vec)
- Binarizing (e.g. needed for Deep Learning multi-class target vector encoding)
- Encoding strings as int
- Dimensionality Reduction (e.g. PCA)
- Augmentation (e.g. tild/zoom images)
- Feature selection based on classifier
- Variance threshold

## Data Types



01

Primary Python libraries: `pandas`, `sklearn`, `scipy`





# Imputing Missing Values

```
# using the most_frequent value
df.src_bytes = df['src_bytes'].fillna
(df['src_bytes'].value_counts().index[0])

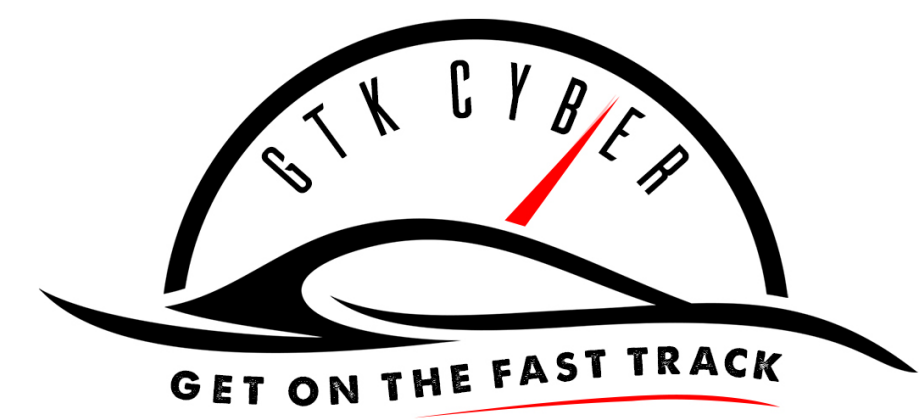
# using the mean value
df.dst_bytes = df['dst_bytes'].fillna(df['dst_bytes'].mean())
```






# One Hot Encoding

Color
Red
Red
Blue
Green
Yellow
Red



# One Hot Encoding

4 Categories



4 Columns with 1 when Category is True and delete original column!

Color
Red
Red
Blue
Green
Yellow
Red

Color_Red	Color_Blue	Color_Yellow	Color_Green
1	0	0	0
1	0	0	0
0	1	0	0
0	0	0	1
0	0	1	0
1	0	0	0



# One Hot Encoding

```
colors = ['Red', 'Red', 'Blue', 'Green', 'Yellow', 'Red']  
series_data = pd.Series( colors )  
pd.get_dummies( series_data )
```

```
# df scenario  
df=pd.get_dummies(df, prefix=None, prefix_sep='_',  
dummy_na=False, columns=['protocol_type','flag'],  
sparse=False)
```



# Encoding Strings as Integers

```
PROBE = ['portsweep.', 'satan.', 'nmap.', 'ipsweep.']  
df = df.replace(to_replace = PROBE, value=1)
```



# Feature Scaling





When you're creating a scaling object, you should first "**fit**" it to the **training data**, then **transform** both the **training and testing data** using the "fit" scaler.

If you try to fit the training and testing data separately, **you will get inaccurate results.**



# Standard Scaling

$$zscore(x_i) = \frac{x_i - \mu}{\sigma}$$

`scaled_feature = (feature - column_mean) / standard_deviation`



# Standard Scaling

$$zscore(x_i) = \frac{x_i - \mu}{\sigma}$$

```
from sklearn.preprocessing import StandardScaler
```

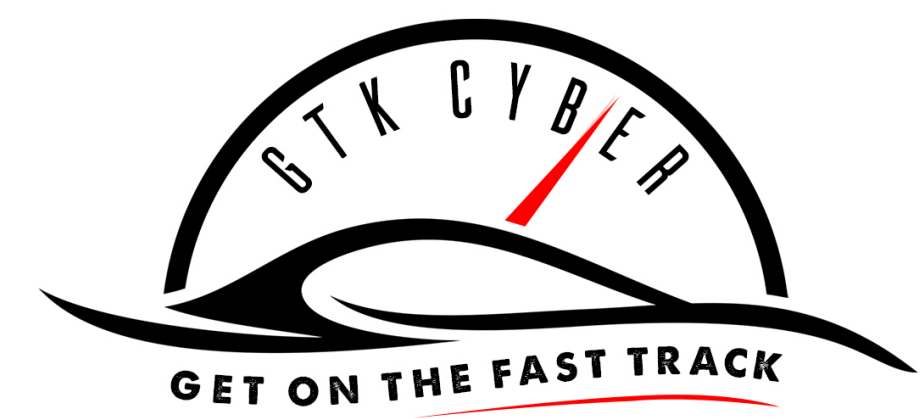
```
scaler = StandardScaler()
```

```
scaler.fit(features)
```

```
features_scaled = scaler.transform(features)
```

or

```
features_scaled = scaler.fit_transform(features)
```



# Standard Scaling

original values:

```
[[ 0.9  0.1 40. ]  
 [ 0.3  0.2 50. ]  
 [ 0.6  0.8 60. ]]
```

Mean of each column:

```
[ 0.6  0.3667 50.]
```

SD of each column:

```
[ 0.2449  0.3091 8.165 ]
```

scaled values:

```
[[ 1.2247 -0.8627 -1.2247]  
 [-1.2247 -0.5392  0.      ]  
 [ 0.      1.4018  1.2247]]
```

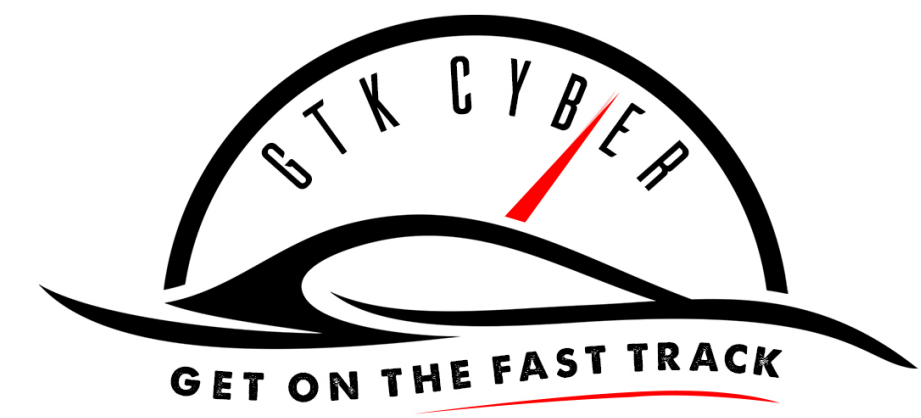
**Means of scaled data, per column:**

```
[ 0. -0.  0.]
```

**SD's of scaled data, per column:**

```
[ 1.  1.  1.]
```

Notice that the mean of standard scaled data is zero and the StdDev is 1.



# Min/Max Scaling

$$X_{norm} = \frac{X - X_{min}}{X_{max} - X_{min}}$$

```
normed_feature = (feature - col_min) / (col_max - col_min)
```





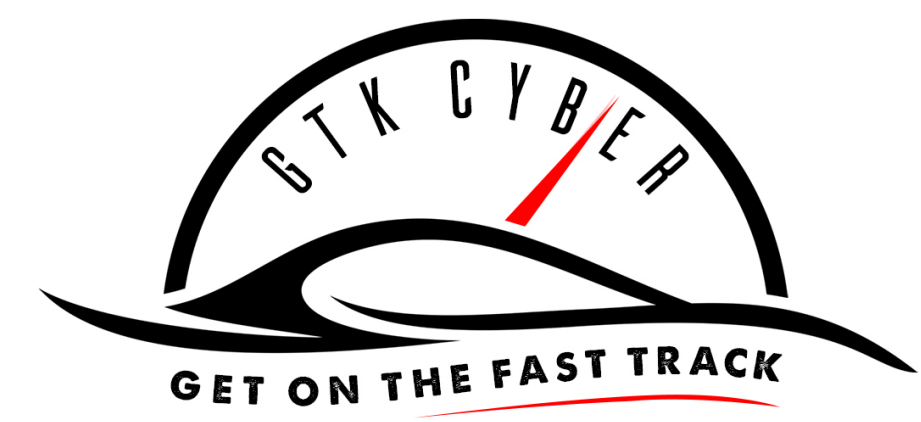
# Min/Max Scaling

```
from sklearn.preprocessing import MinMaxScaler  
minmax = MinMaxScaler()
```

```
minmax.fit(features)  
features_scaled_minmax = minmax.transform(features)
```

or

```
features_scaled_minmax = minmax.fit_transform(features)
```



# Min/Max Scaling

original values:

```
[[ 0.9  0.1 40. ]  
 [ 0.3  0.2 50. ]  
 [ 0.6  0.8 60. ]]
```

Mean of each column:

```
[ 0.6  0.3667 50.]
```

SD of each column:

```
[ 0.2449  0.3091 8.165 ]
```

scaled values:

```
[[ 1.  0.  0. ]  
 [ 0.  0.1429 0.5 ]  
 [ 0.5  1.  1. ]]
```


Means of scaled data, per column:

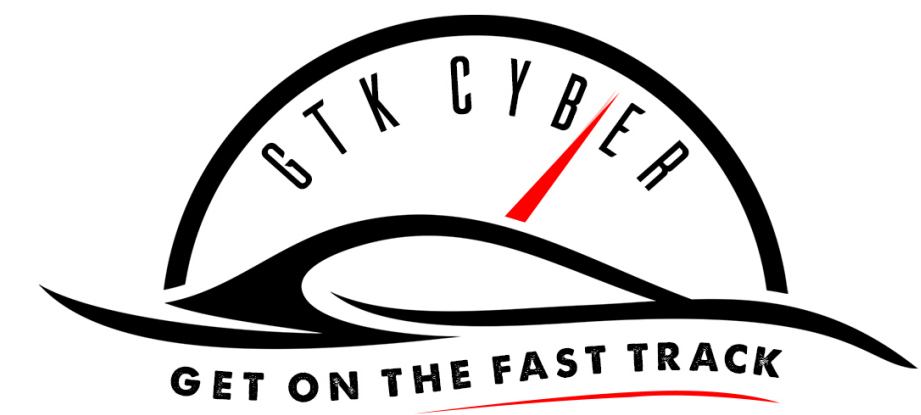
```
[ 0.5  0.381 0.5 ]
```

SD's of scaled data, per column:

```
[ 0.4082  0.4416 0.4082]
```







# Feature Tools

- Open Source
- Automated Feature Engineering

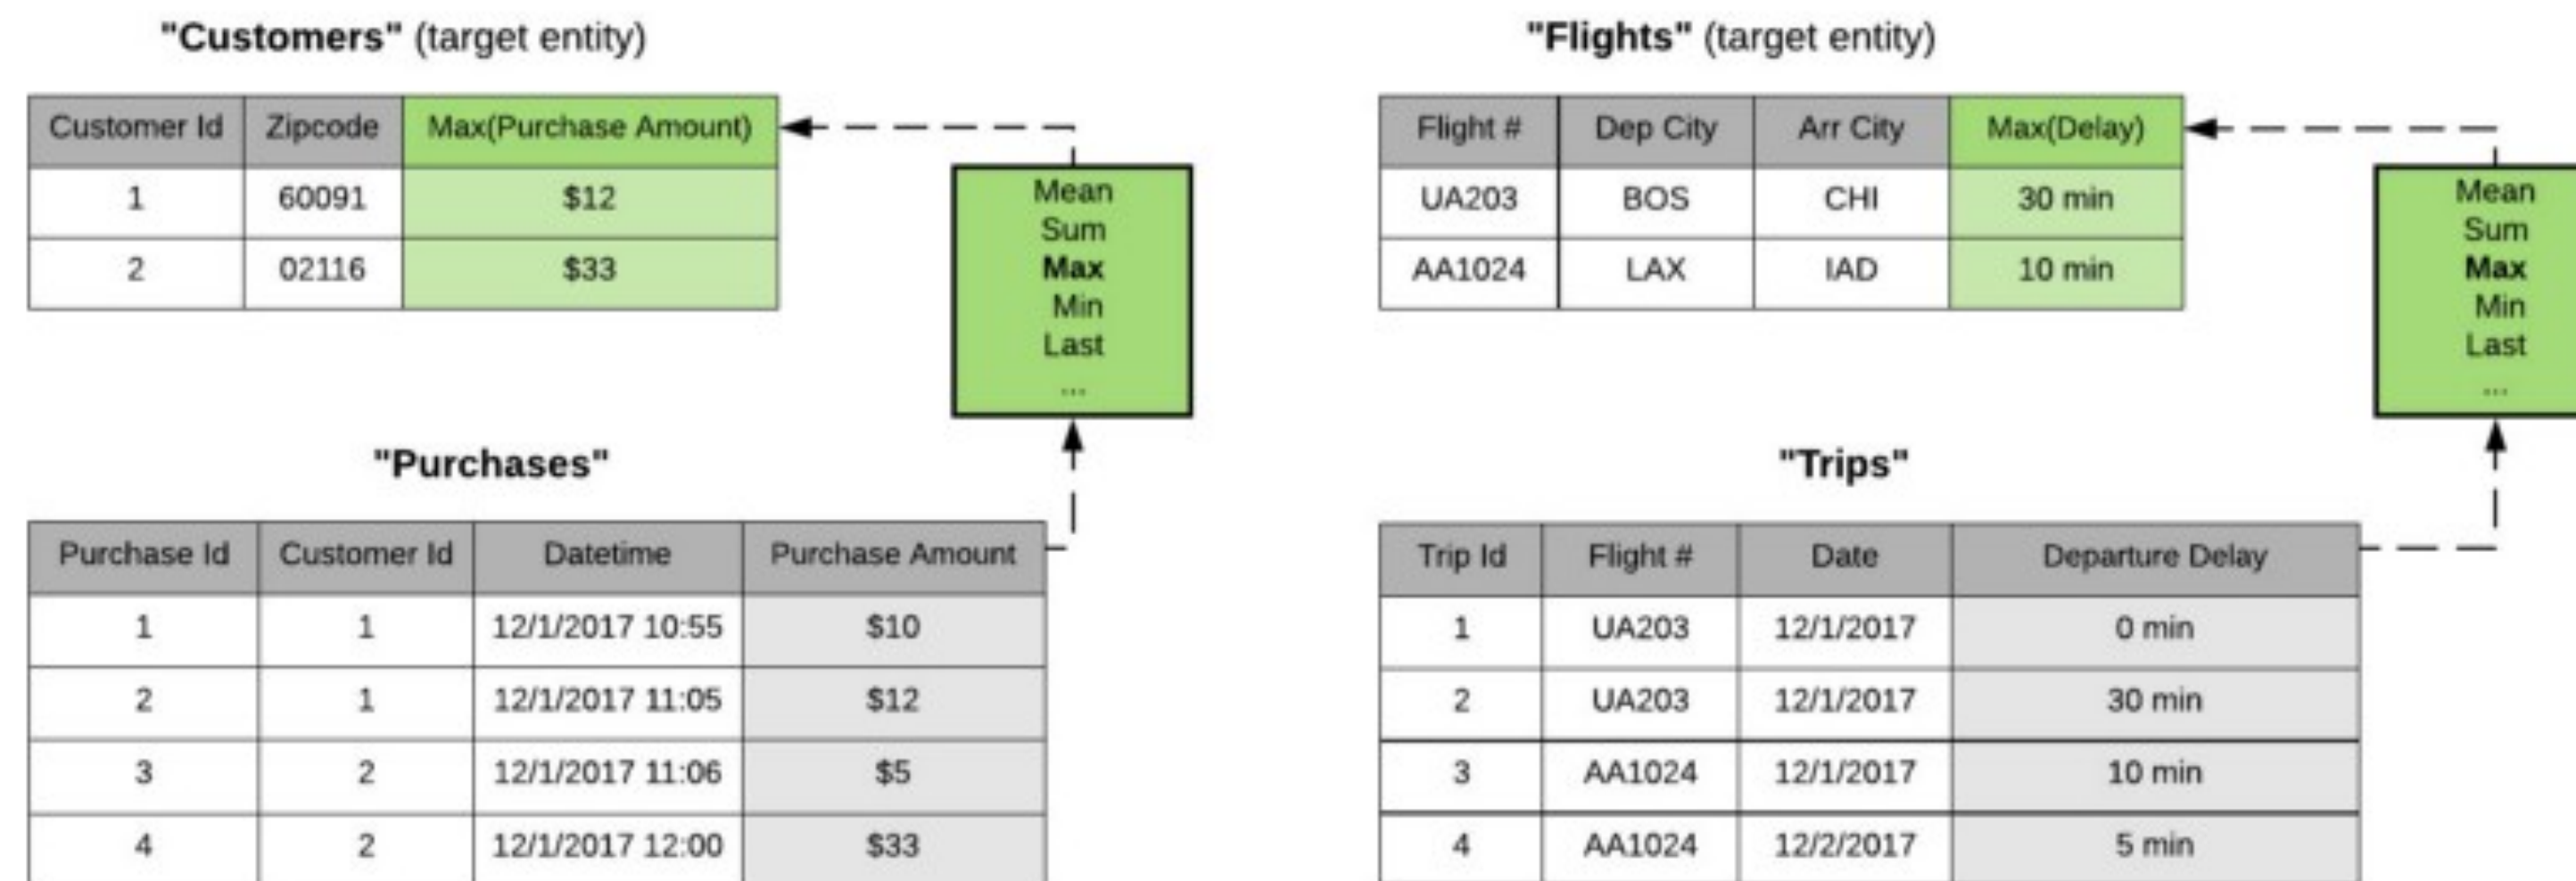


<https://www.featuretools.com/>



# Feature Tools

1. Features are derived from relationships between the data points in a dataset.




To calculate a customer's most expensive purchase, we apply the **Max** primitive to the purchase amount field in all related purchases. When we perform the same steps to a dataset of airplane flights, we calculate "the longest flight delay".

<https://www.featuretools.com/>

# Feature Tools

2. Across datasets, many features are derived by using similar mathematical operations.



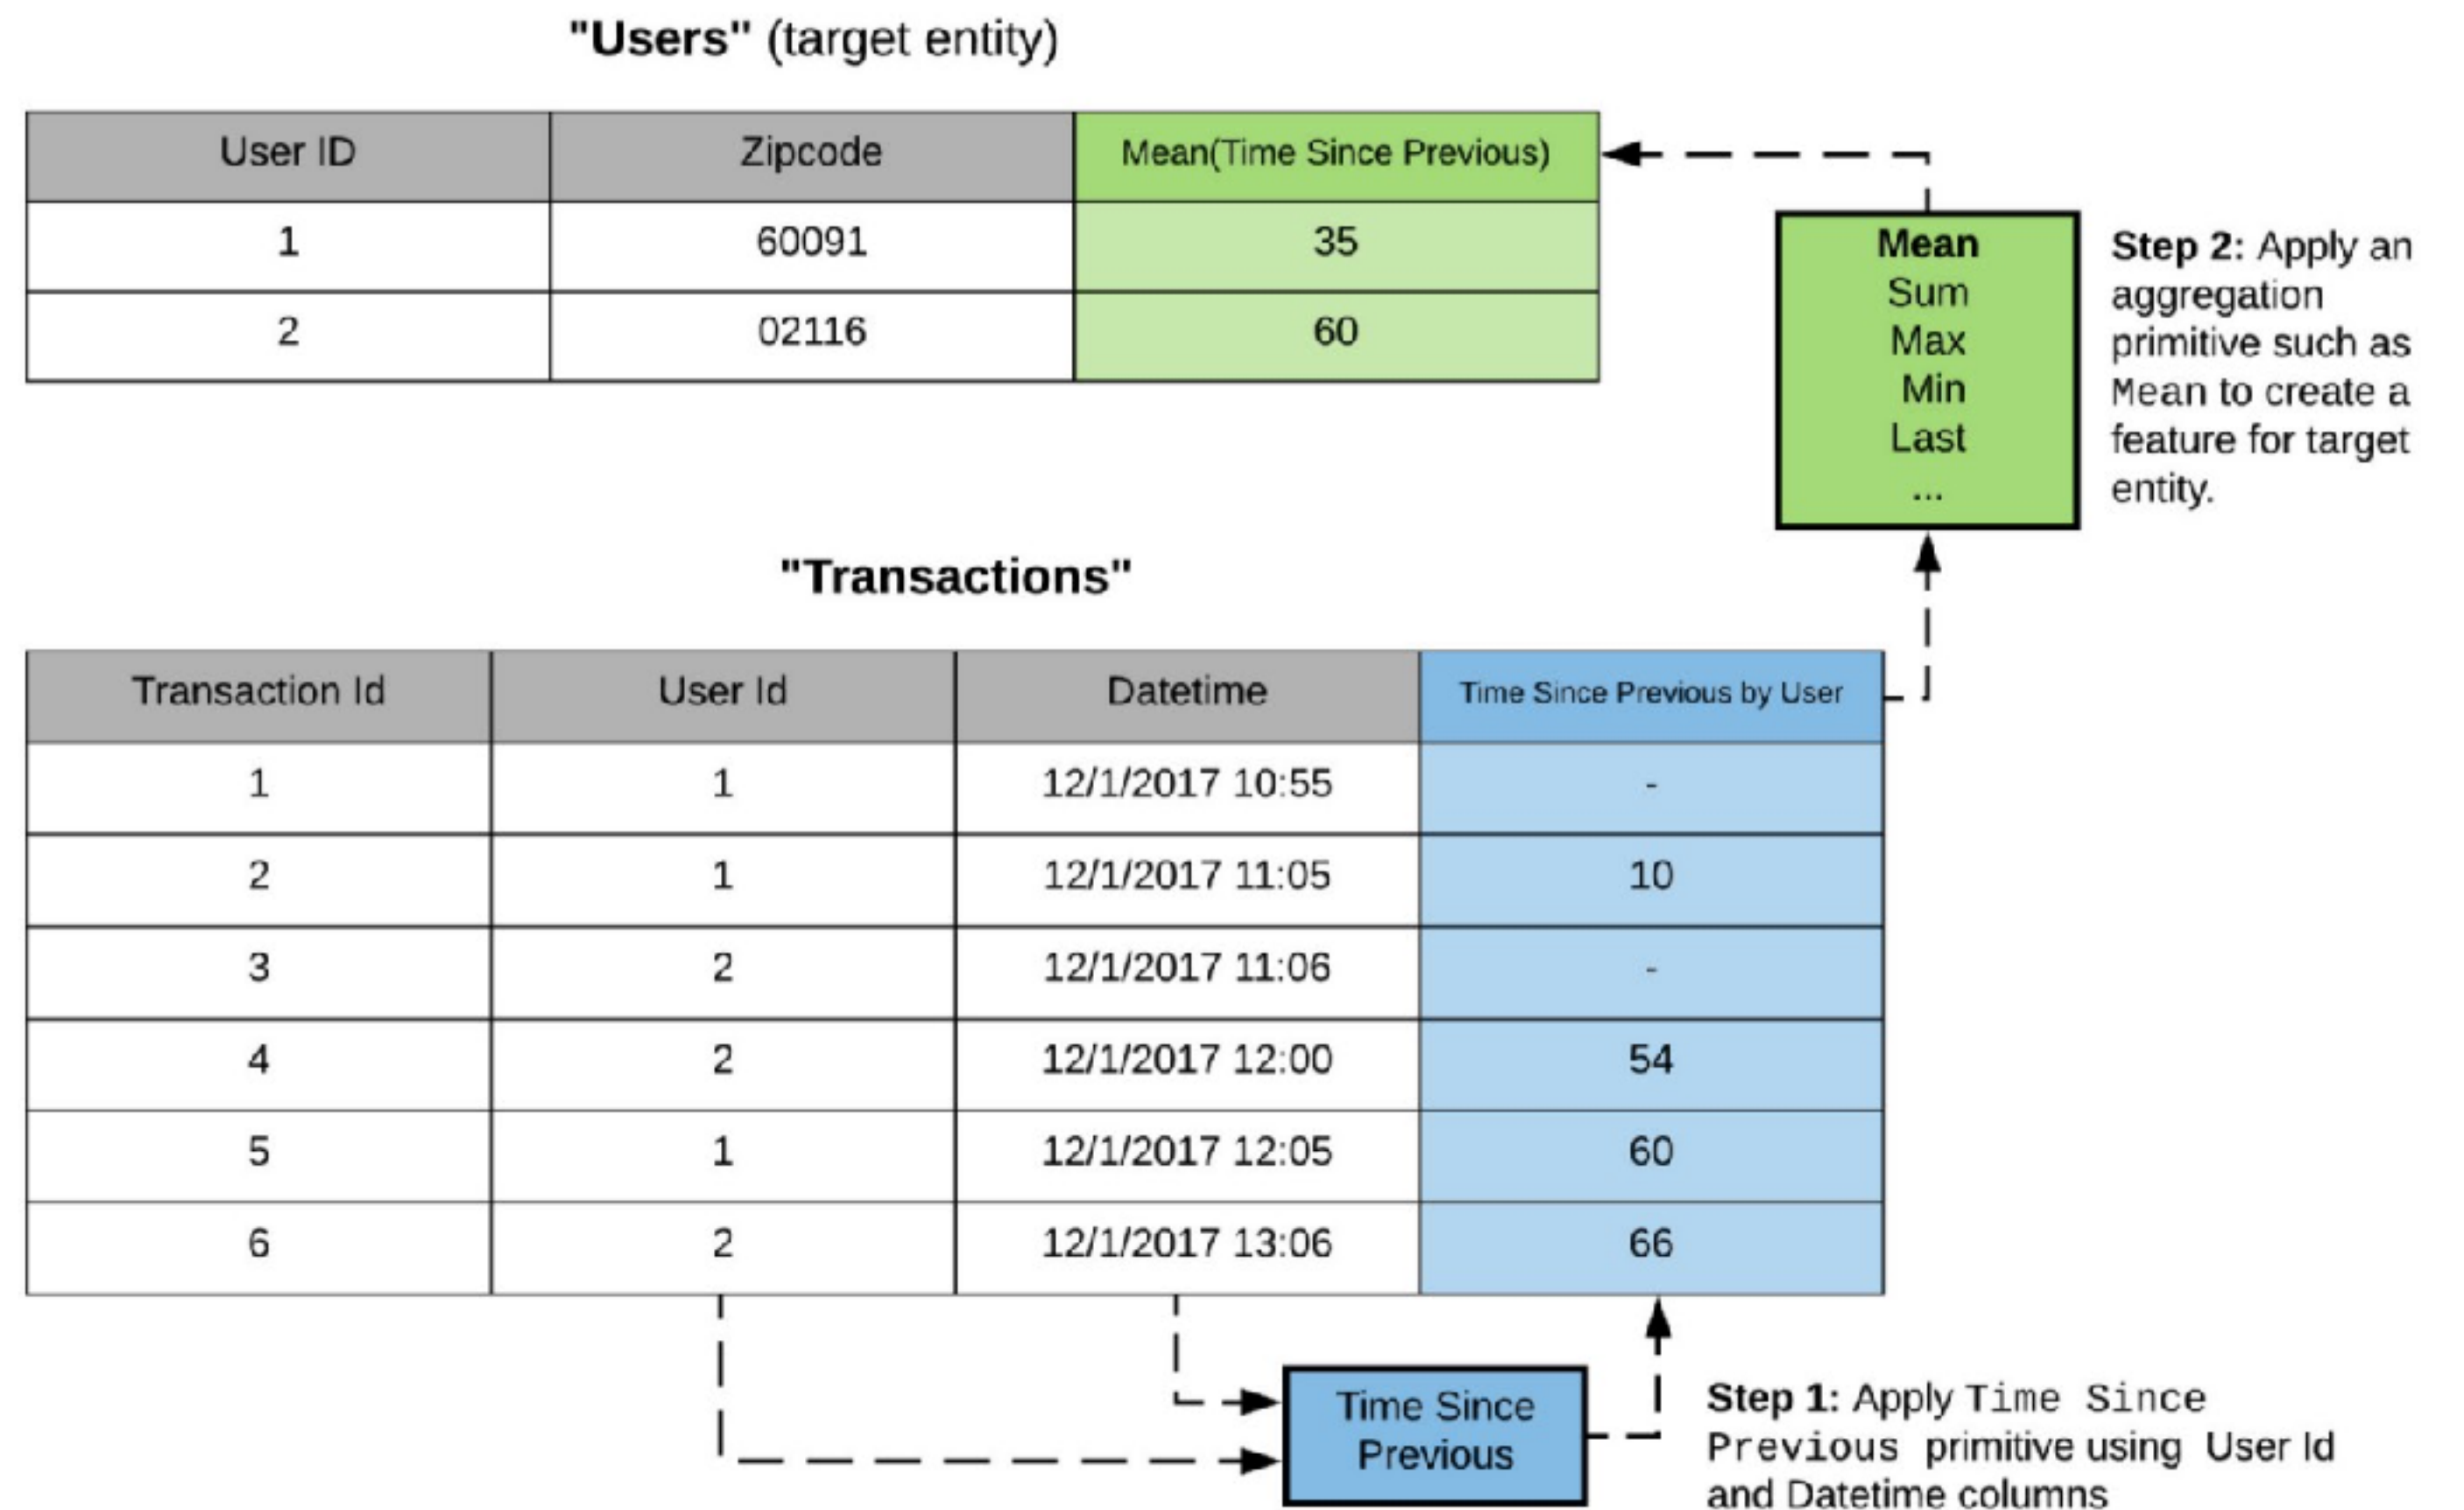
To calculate a customer's most expensive purchase, we apply the **Max** primitive to the purchase amount field in all related purchases. When we perform the same steps to a dataset of airplane flights, we calculate "the longest flight delay".

<https://www.featuretools.com/>




# Feature Tools

3. New features are often composed from utilizing previously derived features.



<https://www.featuretools.com/>




# Selecting Features



Should we use all of them?



How do we know which features to use  
and which to discard?



Selects  $k$  features according to the highest score

```
best_features = SelectKBest(score_func=chi2,k=3).fit_transform(features,target)
```

Selects all features above a given threshold in the scoring function

```
best_features =  
SelectPercentile(score_func=chi2,percentile=3).fit_transform(features,target)
```

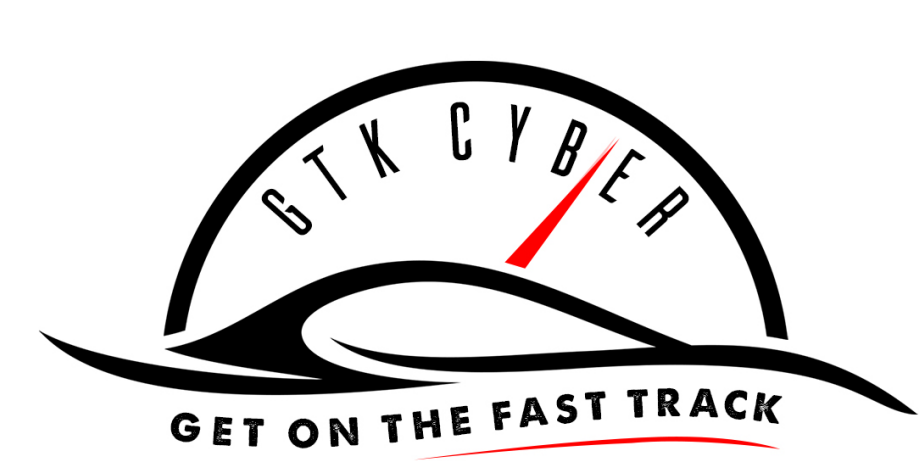
Available Scoring Functions:

- **For regression:** `f_regression`, [`mutual\_info\_regression`](#)
- **For classification:** `chi2`, `f_classif`, `mutual_info_classif`

*References:*

[http://scikit-learn.org/stable/modules/generated/sklearn.feature\\_selection.SelectKBest.html](http://scikit-learn.org/stable/modules/generated/sklearn.feature_selection.SelectKBest.html)

[http://scikit-learn.org/stable/modules/feature\\_selection.html#univariate-feature-selection](http://scikit-learn.org/stable/modules/feature_selection.html#univariate-feature-selection)




How do we know which features to use  
and which to discard?


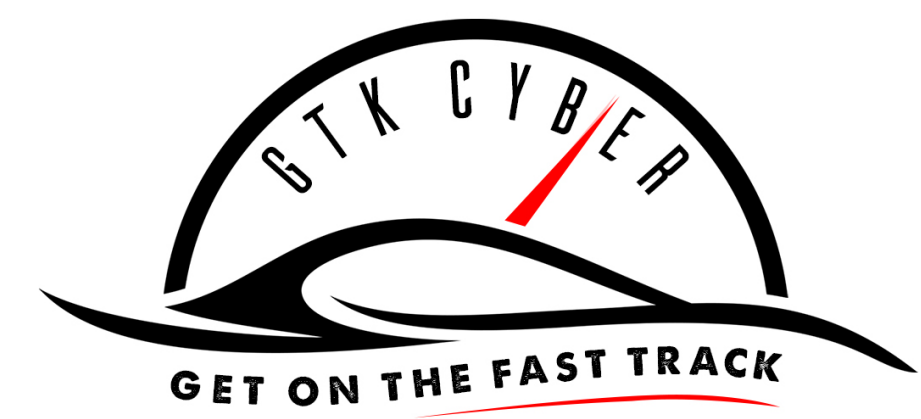
Visualize them!!





# Introducing Yellowbrick and scikit-plot






```
import pandas as pd
import seaborn as sns
from yellowbrick.features.rankd import Rank2D


# Extract the numpy arrays from the data frame
features = df[features]
target = df.diagnosis

# Instantiate the visualizer with the Covariance ranking algorithm
visualizer = Rank2D(features=features, algorithm='covariance')

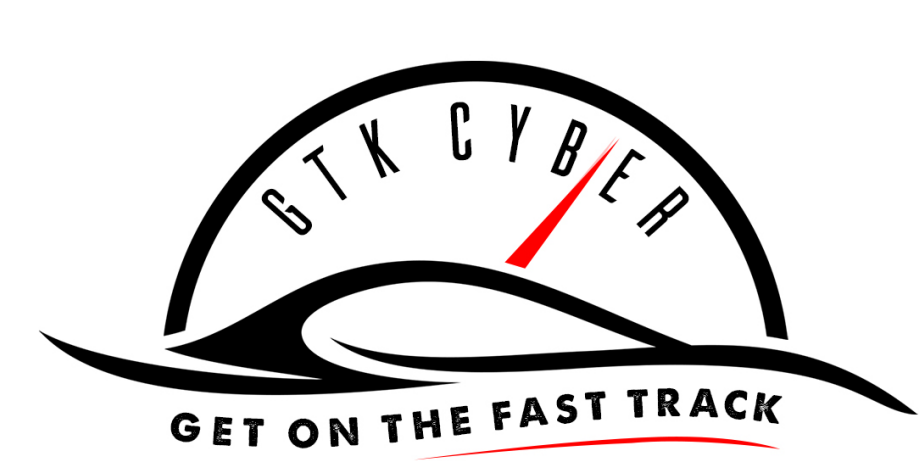
visualizer.fit(features, target)
visualizer.transform(features)
visualizer.poof()
```



```
import seaborn as sns
sns.pairplot(<features>, hue='<target>' )
```



```
from yellowbrick.features.radviz import RadViz
...
visualizer = RadViz(classes=<target classes>, features = <features>)
visualizer.fit(features, target)
visualizer.transform(features)
visualizer.poof()
```



# In Class Exercise

Please take 45 minutes and complete  
**Worksheet 5.1 - Feature Engineering**